

Santander Private Banking UK requires a case study and model

1) To Identify and visualize which factors contribute to customer churn

2) Build a prediction model that will perform the following:

Classify if a customer is going to churn or not.

Preferably and based on model performance, choose a model that will attach a probability to the churn to make it easier for customer service to target low hanging fruits in their efforts to prevent churn.

```
1 cd /content/drive/MyDrive/Colab_Notebooks/DL/Krish_Naik/Churn_Modelling
/content/drive/MyDrive/Colab_Notebooks/DL/Krish_Naik/Churn_Modelling

1 #! pip list
2

1 !pip install tensorflow
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.12.0)
Requirement already satisfied: tensorboard<2.13,>=2.12 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: keras<2.13,>=2.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.
Requirement already satisfied: jax>=0.3.15 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.8)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.3
Requirement already satisfied: tensorflow-estimator<2.13,>=2.12.0 in /usr/local/lib/python3.10/dist-packages (fro
Requirement already satisfied: numpy<1.24,>=1.22 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.8.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: flatbuffers>=2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (23.
Requirement already satisfied: gast<0.4.0,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow) (67.7.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (23.1)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorfl
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 i
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (16.
Requirement already satisfied: wrapt<1.15,>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.
Requirement already satisfied: scipy>=1.7 in /usr/local/lib/python3.10/dist-packages (from jax>=0.3.15->tensoflo
Requirement already satisfied: ml-dtypes>=0.0.3 in /usr/local/lib/python3.10/dist-packages (from jax>=0.3.15->ten
Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in /usr/local/lib/python3.10/dist-packages (from te
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.13,
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages (from tensorboard
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.13,
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from ten
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from google-auth
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from google-aut
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from google-a
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
Requirement already satisfied: charset-normalizer~2.0.0 in /usr/local/lib/python3.10/dist-packages (from request
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1
Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in /usr/local/lib/python3.10/dist-packages (from pyasn1-modul
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages (from requests-oauthlib
```

```

1 #!pip install --upgrade tensorflow
2
3 # Tensorflow version should be greater than 2.0.0 , so that Keras will be supported in version higher than 2.0

```

```

1 import tensorflow as tf
2 tf.__version__

'2.12.0'

```

```

1 #import some basic libraries
2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import pandas as pd
6

```

```

1 dataset=pd.read_csv('Churn_Modelling.csv')
2 dataset.head()

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	151600
1	2	15647311	Hill	608	Spain	Female	41	1	86500
2	3	15619304	Onio	502	France	Female	42	8	159660
3	4	15701354	Boni	699	France	Female	39	1	59180
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510

```

1 # Divide the Dataset into Independent and Dependent features
2
3 # Row number , Customer Id & Surname is not required, so we are removing
4
5 X = dataset.iloc[:,3:13] # from 3rd column to 13th column
6 y = dataset.iloc[:,13]
7

```

```
1 X.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	619	France	Female	42	2	0.00	1	1
1	608	Spain	Female	41	1	83807.86	1	0
2	502	France	Female	42	8	159660.80	3	1
3	699	France	Female	39	1	0.00	2	0
4	850	Spain	Female	43	2	125510.82	1	1

```
1 y.head()
```

```

0    1
1    0
2    1
3    0
4    0
Name: Exited, dtype: int64

```

Feature Engineering

```
1 pd.get_dummies(X['Geography'])
```

	France	Germany	Spain
0	1	0	0
1	0	0	1
2	1	0	0
3	1	0	0
4	0	0	1
...
9995	1	0	0
9996	1	0	0
9997	1	0	0
9998	0	1	0
9999	1	0	0

10000 rows × 3 columns

```
1 geography = pd.get_dummies(X['Geography'], drop_first=True)
2 geography.head(5)
```

	Germany	Spain
0	0	0
1	0	1
2	0	0
3	0	0
4	0	1

```
1 gender = pd.get_dummies(X['Gender'], drop_first=True)
2 gender.head(5)
```

	Male
0	0
1	0
2	0
3	0
4	0

```
1 ## Concatenate these variables with dataframe
2
3 X = X.drop(['Geography', 'Gender'], axis=1) # axis= 1 will remove only the column.
```

```
1 X = pd.concat([X, geography, gender], axis=1)
```

▼ Splitting the dataset into Training & Test set

```
1 from sklearn.model_selection import train_test_split
2
```

```
3 X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state=0)
```

▼ Feature Scaling:

```
1 from sklearn.preprocessing import StandardScaler
2 sc = StandardScaler()
3
4 X_train = sc.fit_transform(X_train)
5 X_test = sc.transform(X_test)
```

```
1 X_train
```

```
array([[ 0.16958176, -0.46460796,  0.00666099, ..., -0.5698444 ,
         1.74309049, -1.09168714],
       [-2.30455945,  0.30102557, -1.37744033, ...,  1.75486502,
        -0.57369368,  0.91601335],
       [-1.19119591, -0.94312892, -1.031415  , ..., -0.5698444 ,
        -0.57369368, -1.09168714],
       ...,
       [ 0.9015152 , -0.36890377,  0.00666099, ..., -0.5698444 ,
        -0.57369368,  0.91601335],
       [-0.62420521, -0.08179119,  1.39076231, ..., -0.5698444 ,
         1.74309049, -1.09168714],
       [-0.28401079,  0.87525072, -1.37744033, ...,  1.75486502,
        -0.57369368, -1.09168714]])
```

```
1 X_test
```

```
array([[ -0.55204276, -0.36890377,  1.04473698, ...,  1.75486502,
        -0.57369368, -1.09168714],
       [-1.31490297,  0.10961719, -1.031415  , ..., -0.5698444 ,
        -0.57369368, -1.09168714],
       [ 0.57162971,  0.30102557,  1.04473698, ..., -0.5698444 ,
         1.74309049, -1.09168714],
       ...,
       [-0.74791227, -0.27319958, -1.37744033, ..., -0.5698444 ,
         1.74309049,  0.91601335],
       [-0.00566991, -0.46460796, -0.33936434, ...,  1.75486502,
        -0.57369368,  0.91601335],
       [-0.79945688, -0.84742473,  1.04473698, ...,  1.75486502,
        -0.57369368,  0.91601335]])
```

```
1 X_train.shape
```

```
(8000, 11)
```

▼ Lets create ANN

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Dense
3 from tensorflow.keras.layers import LeakyReLU, PReLU, ELU, ReLU #Activation functions
4 from tensorflow.keras.layers import Dropout
5
```

▼ Initialize ANN

```
1 classifier = Sequential()
```

```

1 # Adding input layer
2
3 classifier.add(Dense(units=11, activation='relu')) # Dense is used for adding input layer, Hidden layer & O/P layer
4 # units = 11 , number of input columns
5 # Activation will be Relu. This will be added to the next layer

1 # Adding 1st Hidden layer with 7 Neurons
2
3 classifier.add(Dense(units=7, activation='relu')) # Dense is used for adding input layer, Hidden layer & O/P layer
4 classifier.add(Dropout(0.2)) # units = 7 , number of neurons in 1st hidden layer
5 # Activation will be Relu.
6 # Dropout will deactivate the neuron once the value of 0.2 is r

1 # Adding 2nd Hidden layer with 6 Neurons
2
3 classifier.add(Dense(units=6, activation='relu')) # Dense is used for adding input layer, Hidden layer & O/P layer
4 # units = 6 , number of neurons in 2nd hidden layer
5 # Activation will be Relu. This will be added to the next layer

1 # Adding the O/P layer
2
3 classifier.add(Dense(units=1, activation='sigmoid')) # # Dense is used for adding input layer, Hidden layer & O/P
4 # units = 1 , number of neurons in O/P layer
5 # Activation will be Sigmoid. Since its a Binary Classificati

```

▼ Train my ANN

```

1 classifier.compile(optimizer='Adam', loss='binary_crossentropy', metrics=['accuracy']) # Adam the best Optimizer. /
2 # Loss Function will be 'Bir
3 # Metrics will be 'Accurcay'

1 # For information purpose
2 #In case if you want to add your own learning rate, you can import it from the below library
3
4 #import tensorflow
5 #opt = tensorflow.keras.optimizers.Adam(learning_rate=0.01)
6
7 # Then add the optimizer values like below
8 # classifier.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

```

▼ Training our Neural Network

```

1 #Early stopping: when the Training Accuracy is not increasing for a long time, at that time, we can stop(Early Stop)
2
3 #Google 'Early Stopping in Keras'.
4
5 #tf.keras.callbacks.EarlyStopping(
6 #     monitor="val_loss",
7 #     min_delta=0,
8 #     patience=0,
9 #     verbose=0,
10 #     mode="auto",
11 #     baseline=None,
12 #     restore_best_weights=False,
13 #     start_from_epoch=0,
14 # )
15
16
17 #Early Stopping; You can play around with the values

```

```
18
19 import tensorflow as tf
20 early_stopping = tf.keras.callbacks.EarlyStopping(
21     monitor="val_loss",
22     min_delta=0.0001,
23     patience=20,
24     verbose=1,      # Verbose should be 1 as I need to see all the details
25     mode="auto",
26     baseline=None,
27     restore_best_weights=False,
28     start_from_epoch=0,
29 )

1 model_history = classifier.fit(X_train, y_train, validation_split=0.33, batch_size=10, epochs=1000, callbacks=early_
2
3 # Validation_split =.33 ==> out of 100% of training data, we take (100 - 33 = 66) 66% of data for validation.
4
5 # 536/536 is the iteration based on the epoch and Batch size
6
7 # loss => training loss
8 # val_loss = validation loss
9
10 # After some point of time my accuray will be stagnant, at that point of time,
11
12 #Since in the above iteration, most of the time, its going around 88%. so I am going to stop . Stop the google colab
13 #introduce Early stopping
14
15 # early_stopping ==> though we gave epochs = 10000 , it will not run the epoch until 1000 as its a time consuming. c
16 # and if the Accuracy value is not getting increased after a particular time, the call_back will be executed and the
17 # If the Val_loss is not improving much , it will be stopped.
18
```

```

536/536 [=====] - 2s 4ms/step - loss: 0.3269 - accuracy: 0.8666 - val_loss: 0.3595 - v
Epoch 101/1000
536/536 [=====] - 3s 6ms/step - loss: 0.3297 - accuracy: 0.8638 - val_loss: 0.3593 - v
Epoch 102/1000
536/536 [=====] - 3s 6ms/step - loss: 0.3280 - accuracy: 0.8645 - val_loss: 0.3591 - v
Epoch 103/1000
536/536 [=====] - 2s 4ms/step - loss: 0.3279 - accuracy: 0.8642 - val_loss: 0.3598 - v
Epoch 104/1000
536/536 [=====] - 2s 4ms/step - loss: 0.3257 - accuracy: 0.8660 - val_loss: 0.3596 - v
Epoch 105/1000
536/536 [=====] - 3s 5ms/step - loss: 0.3256 - accuracy: 0.8628 - val_loss: 0.3585 - v
Epoch 106/1000
536/536 [=====] - 3s 6ms/step - loss: 0.3274 - accuracy: 0.8653 - val_loss: 0.3591 - v
Epoch 107/1000
536/536 [=====] - 3s 5ms/step - loss: 0.3250 - accuracy: 0.8668 - val_loss: 0.3616 - v
Epoch 108/1000
536/536 [=====] - 2s 4ms/step - loss: 0.3260 - accuracy: 0.8675 - val_loss: 0.3615 - v
Epoch 109/1000

```

In the above O/P, its "Epoch 57: early stopping" with Accuracy 0.8707

Double-click (or enter) to edit

```

1 model_history.history.keys()

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])


```

Summary History for Accuracy

```


1 plt.plot(model_history.history['accuracy'])
2 plt.plot(model_history.history['val_accuracy'])
3 plt.title('Model Accuracy')
4 plt.ylabel('accuracy')
5 plt.xlabel('epoch')
6 plt.legend(['train', 'test'], loc='upper left')
7 plt.show()
8
9 # IN the below diagram, the size/gap is almost same or very less.

```



Summary History for Loss

```
1 plt.plot(model_history.history['loss'])
2 plt.plot(model_history.history['val_loss'])
3 plt.title('Model loss')
4 plt.ylabel('loss')
5 plt.xlabel('epoch')
6 plt.legend(['train', 'test'], loc='upper left')
7 plt.show()
```



Making the Predictions and evaluating the Model

```
1 # Predicting Test set results
2 y_pred = classifier.predict(X_test)
3 y_pred = (y_pred >= 0.5) # If its GEQ 0.5 I will take it as 1 else 0
4
```

63/63 [=====] - 0s 2ms/step

Create Confusion Matrix

```
1 from sklearn.metrics import confusion_matrix
2 cm = confusion_matrix(y_test, y_pred)
3 cm
```

```
array([[1507,  88],
       [ 201, 204]])
```

To calculate Accuracy


```
1 from sklearn.metrics import accuracy_score
2 score=accuracy_score(y_pred,y_test)
3 print('Accuracy :', score)
```

Accuracy : 0.8555

▼ To get the weights assigned to Neural Network

```
1 classifier.get_weights() # These are all the weights value

0.07202806, 0.5394073, 0.33068386, -0.23740686, -0.49479774,
0.34787634],
[-0.3795275, 0.547824, -0.26878467, 0.24962221, -0.6195596,
-0.68806213, 0.31104556, 0.00565256, 0.0132001, 0.16133149,
-0.13530122],
[0.3740936, 0.33844286, -0.03629395, 0.15874322, 0.23995042,
0.31639075, 0.01470504, 0.10974654, 0.04500822, -0.03579355,
-0.04594004]], dtype=float32),
array([0.50589293, -0.08963937, 0.91366386, 0.09575628, -0.36327595,
-0.13635701, -0.03612503, -0.10989795, -0.95137244, -0.47099164,
-0.02579247], dtype=float32),
array([[0.0166559, -0.17833902, 0.06943022, 0.5816181, -0.5013056,
0.5250197, 0.0736379],
[-0.02585014, -0.01116616, 0.4164095, 0.5307345, -0.0553316,
0.4187833, 0.30748808],
[-1.2726372, -0.6833254, 0.08398324, 0.40602022, -0.48742688,
0.6639144, 0.26270002],
[-0.23310532, -0.4241048, 0.08924681, 0.29747796, -0.57601947,
0.18386602, 0.31719837],
[0.46666396, 0.45733413, 0.36872295, -0.5179982, 0.38955614,
-0.5727541, -0.11012888],
[0.13256729, -0.7387074, 0.24270202, 0.30282652, -0.39864323,
0.40435734, 0.16977046],
[0.52163255, 0.17678253, -0.05061479, -0.20306112, 0.2721399,
-0.10630787, -0.8413602],
[-0.43188822, 0.188136, -1.368699, -0.6481819, -0.02907252,
0.06124942, 0.3365523],
[1.2982361, 0.9704765, -0.41823176, -1.051541, 0.9643132,
-1.0170395, -0.4345287],
[-0.60650176, -0.51276344, -0.5057549, -0.0485401, -0.02069807,
0.15953368, -0.21401075],
[0.53626853, 0.21261051, -0.26390353, -0.8310787, 0.24866629,
-0.7560894, -0.8292469]], dtype=float32),
array([-0.04860691, -0.06202723, 0.3498477, 0.8805351, 0.09045316,
0.16145904, -0.15595785], dtype=float32),
array([[0.03264572, 0.9465647, 0.10119472, -0.3636538, -0.7818476,
-0.7780716],
[-0.15745986, 0.6498983, -0.41560486, -1.4087316, -1.1607507,
-1.036178],
[0.81111526, 0.1849275, 0.16152415, 0.54406846, 0.15789898,
0.5996481],
[-0.9223942, -0.04461425, 0.85996616, 0.49070287, -0.86165726,
-0.4968312],
[-0.02504773, 0.80555785, 0.06402837, -0.470465, -0.3331471,
-0.29270843],
[0.6926857, 0.05040722, 0.60584533, -0.39555728, 0.74151,
-0.81154794],
[0.37341484, 0.36859223, 0.8924846, 0.32440257, 0.13066098,
1.0155251]], dtype=float32),
array([0.04574655, 0.1139995, 0.3558023, 0.21981609, 0.12850721,
0.15664038], dtype=float32),
array([-0.7458111,
[1.0752013,
[-0.6146915,
[-0.9058826,
[-0.56294906,
[-1.8829371]], dtype=float32),
array([0.1041213], dtype=float32)]
```

1

1

1

1

1

1

1

1

1

✓ 0s completed at 8:19 PM ● ✕